
Part 1: Theory of Fine-Tuning

Concept Check (Multiple Choice Questions):

1. Correct Answer: B) It customizes the model for specific tasks or domains.
2. Correct Answer: B) When the model loses its generalization ability after excessive fine-tuning on a specific task.

Application Task:

1. Explanation of Transfer Learning (150–200 words):

Transfer learning is like teaching a professional athlete to excel in a new sport. For example, a basketball player already has strong physical fitness, hand-eye coordination, and teamwork skills. When transitioning to volleyball, they don't start from scratch; instead, they build on their existing skills and focus on learning volleyball-specific techniques like serving and spiking.

Similarly, in machine learning, transfer learning involves taking a pre-trained model (e.g., BERT or GPT) that has already learned general language patterns from vast datasets and fine-tuning it for a specific task (e.g., sentiment analysis or legal document classification). The model retains its general knowledge (e.g., grammar, syntax) and adapts to the new task by learning task-specific patterns (e.g., legal jargon or emotional tone). This approach saves time, computational resources, and data compared to training a model from scratch.

For instance, in healthcare, a pre-trained LLM can be fine-tuned to classify medical records into diagnostic categories. The model leverages its general language understanding and learns to recognize medical terms and patterns, making it highly effective for the task.

2. Example Dataset Structure:

Task: Sentiment Analysis on Movie Reviews

Dataset Structure:

text,label

"I loved the movie! The acting was fantastic.",positive

"The film was boring and poorly made.",negative

"An average movie with some good scenes.",neutral

...

Cleaning Steps:

1. Remove special characters and HTML tags.
2. Convert text to lowercase.
3. Remove stopwords (optional, depending on the task).
4. Balance the dataset to ensure equal representation of labels.

Part 2: Practical Fine-Tuning Session

Hands-On Coding Task:

1. Environment Setup:

```
```python
```

```
Install required libraries
```

```
!pip install transformers datasets accelerate
```

```
Verify GPU availability

import torch

print("GPU Available:", torch.cuda.is_available())

'''
```

## 2. Preprocessing Data:

```
```python

from datasets import load_dataset
from transformers import AutoTokenizer

# Load the IMDB dataset
dataset = load_dataset("imdb")

# Load the tokenizer for distilbert-base-uncased
tokenizer = AutoTokenizer.from_pretrained("distilbert-base-uncased")

# Tokenize the dataset
def preprocess_function(examples):
    return tokenizer(examples["text"], truncation=True, padding="max_length",
max_length=512)

tokenized_dataset = dataset.map(preprocess_function, batched=True)

'''
```

3. Model Training:

```
```python

from transformers import AutoModelForSequenceClassification, TrainingArguments, Trainer
```

```
Load the pre-trained model
```

```
model = AutoModelForSequenceClassification.from_pretrained("distilbert-base-uncased",
num_labels=2)
```

```
Define training arguments
```

```
training_args = TrainingArguments(
 output_dir="./results",
 evaluation_strategy="epoch",
 learning_rate=2e-5,
 per_device_train_batch_size=16,
 per_device_eval_batch_size=16,
 num_train_epochs=3,
 weight_decay=0.01,
 save_strategy="epoch",
 logging_dir="./logs",
)
```

```
Initialize the Trainer
```

```
trainer = Trainer(
 model=model,
 args=training_args,
 train_dataset=tokenized_dataset["train"],
 eval_dataset=tokenized_dataset["test"],
)
```

```
Fine-tune the model
```

```
trainer.train()
```

#### 4. Save and Evaluate:

python

Save the fine-tuned model

```
trainer.save_model("./fine-tuned-distilbert-imdb")
```

Evaluate the model

```
eval_results = trainer.evaluate()
```

```
print("Evaluation Results:", eval_results)
```

---

Reflection:

Key Challenges and Solutions:

1. Challenge: Limited GPU memory caused out-of-memory errors during training.

Solution: Reduced batch size (`per\_device\_train\_batch\_size`) and used gradient accumulation.

2. Challenge: Overfitting on the training dataset.

Solution: Added regularization techniques like weight decay and early stopping.

3. Challenge: Low accuracy on the test set.

Solution: Increased the number of epochs and fine-tuned the learning rate.

Suggestions for Improving Performance:

Hyperparameter Tuning: Experiment with different learning rates, batch sizes, and epochs.

Data Augmentation: Add more diverse examples to the dataset to improve generalization.

Advanced Techniques: Use techniques like learning rate scheduling or mixed-precision training.

Larger Model: If resources allow, fine-tune a larger model like `bert-base-uncased` for better performance.

---

By completing this assignment, you'll gain hands-on experience in fine-tuning LLMs and a deeper understanding of the theory behind it. Let me know if you need further clarification or assistance!